

PROFILE

Ph.D. in Computer Science and Technology at Beijing Institute of Technology (advisor: Prof. Heyan Huang), with a CSC-funded joint-training year at Singapore Management University in 2023–2024 (advisors: Prof. Ee-Peng Lim and Prof. Jing Jiang). Research focuses on dialogue systems and natural language processing. Author of **26** peer-reviewed papers, including **11** as first or corresponding author at venues such as ACL, TOIS and AAAI (**7** CCF-A, **1** SCI Q1, **1** CAAI-A), and **15** co-authored papers at NeurIPS, ACM Computing Surveys and similar venues. Holds 3 granted or filed invention patents and co-authored one monograph. Core contributor to MindLLM, the first foundation model pre-trained end-to-end within a Chinese university, with over 30,000 open-source downloads. Recipient of the First Prize of the Wu Wenjun AI Science & Technology Progress Award (ranked 9th).

EDUCATION

- Oct 2023 – Oct 2024 **Singapore Management University** · School of Computing and Information Systems · Visiting Ph.D. Student (funded by China Scholarship Council)
Advisors: Prof. Ee-Peng Lim (PAKDD Steering Committee life member and Chair) and Prof. Jing Jiang (ACL Vice-President)
- Sep 2020 – Jun 2026 **Beijing Institute of Technology** · School of Computer Science · Ph.D. in Computer Science and Technology (direct M.S.–Ph.D. track)
Advisor: Prof. Heyan Huang (Vice Chair, MOE Teaching Advisory Committee for Computer Science; CAAI and CIPS Fellow). Dissertation: Multi-Dimensional Modeling for Domain-Specific Dialogue 《面向垂域对话的多维度建模方法研究》
- Sep 2016 – Jun 2020 **Beijing Institute of Technology** · School of Computer Science · B.Eng. in Computer Science and Technology
GPA 3.8/4.0 (ranked 9/218); admitted to graduate study by recommendation without examination

PUBLICATIONS

CCF-A SCI Q1 Other tiers * corresponding author

First / Corresponding Author

1. **Yang Y**, Huang H Y, Achananuparp P, Jiang J, Kit P L, Lim N G, Ern C T S, Lim E P. Cami: A counselor agent supporting motivational interviewing through state inference and topic exploration. *ACL 2025 (Long)*: 21037–21081. CCF-A
2. **Yang Y**, Huang H Y, Achananuparp P, Jiang J, Lim N G, Ern C T S, Kit P L, Xiuhui J G, Pinto J, Lim E P. Consistent client simulation for motivational interviewing-based counseling. *ACL 2025 (Long)*: 20959–20998. CCF-A
3. **Yang Y**, Huang H Y, Achananuparp P, Jiang J, Lim E P. Speaker verification in agent-generated conversations. *ACL 2024 (Long)*: 5655–5676. CCF-A
4. **Yang Y**, Huang H, Sun H, Li J, Gao Y. Planner policies are bias learners: A comprehensive analysis of preference bias in LLM-based dialogue planner policies. *ACM Transactions on Information Systems (TOIS)*, 2026. DOI: 10.1145/3831687. CCF-A SCI Q1 · IF 11.2
5. **Yang Y**, Huang H Y, Achananuparp P, Jiang J, Lim E P. MIThinker: A plug-and-play policy-optimized thinker for motivational interviewing counseling. *Findings of ACL 2026 (Long)*: 3292–3328. CCF-A
6. Huang H, **Yang Y***, Sun H, Li J, Gao Y. Simulated rewards, skewed strategies: Tracing the acquired preference bias in LLM-based dialogue planners. *AAAI 2026*, Vol. 40, Issue 26. CCF-A · Oral
7. Xu C, Ji Y, Lv Z, Yi Y, **Yang Y***, Ji Y, Chen C, Wang X, Lan T, Wang Z, et al. PUPPET: Neural-symbolic standardized patients for mental health. *ACL 2026 (Long)*: 2599–2634. CCF-A
8. **Yang Y**, Huang H, Gao Y, Li J. Building knowledge-grounded dialogue systems with graph-based semantic modelling. *Knowledge-Based Systems*, 2024, 298: 111943. SCI Q1 Top · IF 7.6
9. **Yang Y**, Huang H Y, Liu Y, Gao Y. Graph vs. sequence: An empirical study on knowledge forms for knowledge-grounded dialogue. *EMNLP 2023*: 15846–15858. CCF-B / CAAI-A
10. **Yang Y**, Sun H, Li J, Liu R, Li Y, Liu Y, Gao Y, Huang H. MindLLM: Lightweight large language model pre-training, evaluation and domain application. *AI Open*, 2024, 5: 155–180. EI · IF 14.8

11. 杨毅哲, 黄河燕. 基于联律规则约束解码的对联生成方法. *CCL 2026 (The 25th China National Conference on Computational Linguistics)*. CIPS flagship

Co-authored

12. Li J, **Yang Y**, Bai Y, et al. Fundamental capabilities of large language models and their applications in domain scenarios: A survey. *ACL 2024 (Long)*: 11116–11141. CCF-A
13. Li J, Gao Y, **Yang Y**, et al. Fundamental capabilities and applications of large language models: A survey. *ACM Computing Surveys*, 2025, 58(2): 1–42. SCI Q1 Top · IF 28.0
14. Gao Y, Liu Q, **Yang Y**, Wang K. Latent representation discretization for unsupervised text style generation. *Information Processing & Management*, 2024, 61(3): 103643. SCI Q1 Top · ESI Highly Cited
15. Liu Y, Huang H, **Yang Y**, Zeng Z, Zhou Y, Wu Z, Gao Y. Retrieval is NOT always needed: Exploring the timing of retrieval in dynamic retrieval-augmented generation. *Expert Systems with Applications*, 2026, 333(C): 133921. SCI Q1 Top · IF 9.4
16. Tang W, Cao Y, Deng Y, Ying J, Wang B, **Yang Y**, Zhao Y, Zhang Q, Huang X, Jiang Y G, Liao Y. EvoWiki: Evaluating LLMs on evolving knowledge. *ACL 2025 (Long)*: 948–964. CCF-A
17. Ying J, Cao Y, Bai Y, Sun Q, Wang B, Tang W, Ding Z, **Yang Y**, Huang X, Yan S. Automating dataset updates towards reliable and timely evaluation of large language models. *NeurIPS 2024*, 37: 17106–17132. CCF-A
18. Sun H, Wu Y, **Yang Y**, Li Y, Li J, Ye Y, Gao Y. PSST: A benchmark for evaluation-driven text public-speaking style transfer. *Findings of EMNLP 2024*: 8438–8471. CAAI-A
19. Xiao Y, Lin S, Chen C, **Yang Y**, Zhang X, Liao X. Learnable termination safeguard: Token-level early termination for efficient LLM safety monitoring. *EMNLP 2026*. CAAI-A
20. Wang Y, Chen W, Chen C, Liu X, **Yang Y**, Chen Z, He Q, Chen Y, Liu K, Wang X, Liao X, Liu J. ToToFilter: Topic- and topology-aware noise filtering framework over medical KGs for evidence-based medical Q&A. *EMNLP 2026*. CAAI-A
21. Liang X, Li J, **Yang Y**, Gao Y. Bit_numeval at SemEval-2024 Task 7: Enhance numerical sensitivity and reasoning completeness for quantitative understanding. *SemEval-2024*: 1830–1841. 1st place, math reasoning track
22. Liang X, **Yang Y**, Huang H, Gao Y. Auditing large language models as educational assistants and judges. *NLPCC 2026*. CCF-C
23. Lin S, Wu Y, Chen C, **Yang Y**, Chen Z, Yang Z, Liao X. Multi-granularity sentiment integration for LLM-based multimodal sentiment analysis. *NLPCC 2026*. CCF-C · Oral
24. Wang X, Hong D, Chen C, **Yang Y**, Yang Z, Liao X. Co-occurring emotion prototype learning for multi-label emotion recognition in conversation. *NLPCC 2026*. CCF-C
25. Liu Q, Gao Y, **Yang Y**. Unsupervised style transfer in news headlines via discrete style space. *CCL 2023*: 626–647. CIPS flagship
26. Li J, Ren M, Gao Y, **Yang Y**. Ask to understand: Question generation for multi-hop question answering. *CCL 2023*: 569–582. CIPS flagship

PATENTS & MONOGRAPH

- Patent: A multi-granularity text recommendation method based on contextual semantics (ZL 2021 1 0039994.8), second inventor (advisor first), granted 2022
- Patent: A training method for lightweight Chinese–English bilingual large language models (202411485748.5), second inventor (advisor first), filed 2024
- Patent: A two-stage training method for lightweight Chinese–English bilingual large language models (202411485750.2), second inventor (advisor first), filed 2024
- Monograph: *Large Language Models: Technical Practice and Scenario Applications* (ISBN 9787121516139), H. Huang, Z. Chi, Y. Bai, **Y. Yang**, Publishing House of Electronics Industry, 2026

SERVICE & AWARDS

- Reviewer: ICLR, NeurIPS, ICML, ACL, Artificial Intelligence; Program Committee: CIKM, AAI
- Collaborations: National University of Singapore, Nanyang Technological University, Singapore Management University, Australian National University
- 2025 — First Prize, Wu Wenjun AI Science & Technology Progress Award, CAAI “吴文俊人工智能科技进步奖一等奖”(ranked 9th)
- 2025 — National Scholarship for Doctoral Students, Ministry of Education
- 2024 — SemEval-2024 Task 7, mathematical reasoning track, NumEval: Numeral-Aware Language Understanding and Generation

- 2020 — Outstanding Graduate, Beijing Institute of Technology